

## CptS 570 Machine Learning, Fall 2021

### Homework #3

Due Date: Nov 15 midnight via Canvas

NOTE 1: Please use a word processing software (e.g., Microsoft word or Latex) to write your answers and submit a printed copy to me at the beginning of class on Oct 23. The rationale is that it is sometimes hard to read and understand the hand-written answers.

NOTE 2: Please ensure that all the graphs are appropriately labeled (x-axis, y-axis, and each curve). The caption or heading of each graph should be informative and self-contained.

### Analytical Part (6 Percent)

1. Suppose  $x = (x_1, x_2, \dots, x_d)$  and  $z = (z_1, z_2, \dots, z_d)$  be any two points in a high-dimensional space (i.e.,  $d$  is very large).
  - a. Try to prove the following, where the right-hand side quantity represent the standard Euclidean distance.

$$\left( \frac{1}{\sqrt{d}} \sum_{i=1}^d x_i - \frac{1}{\sqrt{d}} \sum_{i=1}^d z_i \right)^2 \leq \sum_{i=1}^d (x_i - z_i)^2 \quad (1)$$

**Hint:** Use Jensen's inequality – If  $X$  is a random variable and  $f$  is a convex function, then  $f(\mathbb{E}[X]) \leq \mathbb{E}[f(X)]$ .

- b. We know that the computation of nearest neighbors is very expensive in the high-dimensional space. Discuss how we can make use of the above property to make the nearest neighbors computation efficient?
2. We briefly discussed in the class about Locality Sensitive Hashing (LSH) algorithm to make the nearest neighbor classifier efficient. Please read the following paper and briefly summarize the key ideas as you understood:  
Alexandr Andoni, Piotr Indyk: Near-optimal hashing algorithms for approximate nearest neighbor in high dimensions. Communications of ACM 51(1): 117-122 (2008) <http://people.csail.mit.edu/indyk/p117-andoni.pdf>
  3. We know that we can convert any decision tree into a set of if-then rules, where there is one rule per leaf node. Suppose you are given a set of rules  $R = \{r_1, r_2, \dots, r_k\}$ , where  $r_i$  corresponds to the  $i^{th}$  rule. Is it possible to convert the rule set  $R$  into an equivalent decision tree? Explain your construction or give a counterexample.
  4. Please read the following paper and briefly summarize the key ideas as you understood (You can skip the proofs, but it is important to understand the main results):

Andrew Y. Ng, Michael I. Jordan: On Discriminative vs. Generative Classifiers: A comparison of logistic regression and naive Bayes. NIPS 2001: 841-848 <http://ai.stanford.edu/~ang/papers/nips01-discriminativegenerative.pdf>

5. Naive Bayes vs. Logistic Regression

- Let us assume that the training data satisfies the Naive Bayes assumption (i.e., features are independent given the class label). As the training data approaches infinity, which classifier will produce better results, Naive Bayes or Logistic Regression? Please explain your reasoning.
- Let us assume that the training data does **NOT** satisfy the Naive Bayes assumption. As the training data approaches infinity, which classifier will produce better results, Naive Bayes or Logistic Regression? Please explain your reasoning.
- Can we compute  $P(X)$  from the learned parameters of a Naive Bayes classifier? Please explain your reasoning.
- Can we compute  $P(X)$  from the learned parameters of a Logistic Regression classifier? Please explain your reasoning.

6. Please read the following paper and briefly summarize the key ideas as you understood:

Thomas G. Dietterich: Ensemble Methods in Machine Learning. Multiple Classifier Systems 2000: 1-15

<https://web.engr.oregonstate.edu/~tgd/publications/mcs-ensembles.pdf>

7. We need to perform statistical tests to compare the performance of two learning algorithms on a given learning task. Please read the following paper and briefly summarize the key ideas as you understood:

Thomas G. Dietterich: Approximate Statistical Test For Comparing Supervised Classification Learning Algorithms. Neural Computation 10(7): 1895-1923 (1998) <http://sci2s.ugr.es/keel/pdf/algorithm/articulo/dietterich1998.pdf>

8. (**Finite-Horizon MDPs.**) Our basic definition of an MDP in class defined the reward function  $R(s)$  to be a function of just the state, which we will call a *state reward function*. It is also common to define a reward function to be a function of the state and action, written as  $R(s, a)$ , which we will call a *state-action reward function*. The meaning is that the agent gets a reward of  $R(s, a)$  when they take action  $a$  in state  $s$ . While this may seem to be a significant difference, it does not fundamentally extend our modeling power, nor does it fundamentally change the algorithms that we have developed.

a) Describe a real world problem where the corresponding MDP is more naturally modeled using a state-action reward function compared to using a state reward function.

b) Modify the Finite-horizon value iteration algorithm so that it works for state-action reward functions. Do this by writing out the new update equation that is used in each iteration and explaining the modification from the equation given in class for state rewards.

c) Any MDP with a state-action reward function can be transformed into an “equivalent” MDP with just a state reward function. Show how any MDP with a state-action reward function  $R(s, a)$  can be transformed into a different MDP with state reward function  $R(s)$ , such that the optimal policies in the new MDP correspond exactly to the optimal policies in the original MDP. That is an optimal policy in the new MDP can be mapped to an optimal policy in the original MDP. *Hint: It will be necessary for the new MDP to introduce new “book keeping” states that are not in the original MDP.*

9. ( **$k$ -th Order MDPs.**) A standard MDP is described by a set of states  $S$ , a set of actions  $A$ , a transition function  $T$ , and a reward function  $R$ . Where  $T(s, a, s')$  gives the probability of transitioning to  $s'$  after taking action  $a$  in state  $s$ , and  $R(s)$  gives the immediate reward of being in state  $s$ .

A  $k$ -order MDP is described in the same way with one exception. The transition function  $T$  depends on the current state  $s$  and also the previous  $k - 1$  states. That is,  $T(s_{k-1}, \dots, s_1, s, a, s') = \Pr(s'|a, s, s_1, \dots, s_{k-1})$  gives the probability of transitioning to state  $s'$  given that action  $a$  was taken in state  $s$  and the previous  $k - 1$  states were  $(s_{k-1}, \dots, s_1)$ .

Given a  $k$ -order MDP  $M = (S, A, T, R)$  describe how to construct a standard (First-order) MDP  $M' = (S', A', T', R')$  that is equivalent to  $M$ . Here equivalent means that a solution to  $M'$  can be easily converted into a solution to  $M$ . Be sure to describe  $S'$ ,  $A'$ ,  $T'$ , and  $R'$ . Give a brief justification for your construction.

10. Some MDP formulations use a reward function  $R(s, a)$  that depends on the action taken in a state or a reward function  $R(s, a, s')$  that also depends on the result state  $s'$  (we get reward  $R(s, a, s')$  when we take action  $a$  in state  $s$  and then transition to  $s'$ ). Write the Bellman optimality equation with discount factor  $\beta$  for each of these two formulations.
11. Consider a trivially simple MDP with two states  $S = \{s_0, s_1\}$  and a single action  $A = \{a\}$ . The reward function is  $R(s_0) = 0$  and  $R(s_1) = 1$ . The transition function is  $T(s_0, a, s_1) = 1$  and  $T(s_1, a, s_1) = 1$ . Note that there is only a single policy  $\pi$  for this MDP that takes action  $a$  in both states.
  - a) Using a discount factor  $\beta = 1$  (i.e. no discounting), write out the linear equations for evaluating the policy and attempt to solve the linear system. What happens and why?
  - b) Repeat the previous question using a discount factor of  $\beta = 0.9$ .

## Programming and Empirical Analysis (2 Percent)

### 1. (40 points) Fortune Cookie Classifier<sup>1</sup>

You will build a binary fortune cookie classifier. This classifier will be used to classify fortune cookie messages into two classes: messages that predict what will happen in the future (class 1) and messages that just contain a wise saying (class 0). For example,

“Never go in against a Sicilian when death is on the line” would be a message in class 0.

“You will get an A in Machine learning class” would be a message in class 1.

**Files Provided** There are three sets of files. All words in these files are lower case and punctuation has been removed.

1) The training data:

traindata.txt: This is the training data consisting of fortune cookie messages.

trainlabels.txt: This file contains the class labels for the training data.

2) The testing data:

testdata.txt: This is the testing data consisting of fortune cookie messages.

testlabels.txt: This file contains the class labels for the testing data.

3) A list of stopwords: stoplist.txt

There are two steps: the pre-processing step and the classification step. In the pre-processing step, you will convert fortune cookie messages into features to be used by your classifier. You will be using a bag of words representation. The following steps outline the process involved:

Form the vocabulary. The vocabulary consists of the set of all the words that are in the training data with stop words removed (stop words are common, uninformative words such as

---

<sup>1</sup>Thanks to Weng-Keen Wong and his advisor Andrew Moore for sharing the data.

“a” and “the” that are listed in the file stoplist.txt). The vocabulary will now be the features of your training data. Keep the vocabulary in alphabetical order to help you with debugging.

Now, convert the training data into a set of features. Let  $M$  be the size of your vocabulary. For each fortune cookie message, you will convert it into a feature vector of size  $M$ . Each slot in that feature vector takes the value of 0 or 1. For these  $M$  slots, if the  $i$ th slot is 1, it means that the  $i$ th word in the vocabulary is present in the fortune cookie message; otherwise, if it is 0, then the  $i$ th word is not present in the message. Most of these feature vector slots will be 0. Since you are keeping the vocabulary in alphabetical order, the first feature will be the first word alphabetically in the vocabulary.

Implement the Naive Bayes Classifier (with Laplace Smoothing) and run it on the training data. Compute the training and testing accuracy.

To debug and test your implementation, you can employ Weka (weka.classifiers.bayes.NaiveBayes): <http://www.cs.waikato.ac.nz/ml/weka/downloading.html> or scikit-learn ([http://scikit-learn.org/stable/modules/naive\\_bayes.html](http://scikit-learn.org/stable/modules/naive_bayes.html))

### Instructions for Code Submission and Output Format.

Please follow the below instructions. It will help us in grading your programming part of the homework. Please submit the zip file on Blackboard.

- Supported programming languages: Python, Java, C++
- Store all the relevant files in a folder and submit the corresponding zipfile named after your student-id, e.g., 114513209.zip
- This folder should have a script file named

`run_code.sh`

Executing this script should do all the necessary steps required for executing the code including compiling, linking, and execution

- Assume relative file paths in your code. Some examples:

`“./filename.txt”` or `“../hw2/filename.txt”`

- The output of your program should be dumped in a file named “output.txt”
- Make sure the output.txt file is dumped when you execute the script

`run_code.sh`

- Zip the entire folder and submit it as

`<student_id>.zip`

## Grading Rubric

Each question in the students work will be assigned a letter grade of either A,B,C,D, or F by the Instructor and TAs. This five-point (discrete) scale is described as follows:

- **A) Exemplary (=100%).**  
Solution presented solves the problem stated correctly and meets all requirements of the problem.  
Solution is clearly presented.  
Assumptions made are reasonable and are explicitly stated in the solution.  
Solution represents an elegant and effective way to solve the problem and is not overly complicated than is necessary.
- **B) Capable (=75%).**  
Solution is mostly correct, satisfying most of the above criteria under the exemplary category, but contains some minor pitfalls, errors/flaws or limitations.
- **C) Needs Improvement (=50%).**  
Solution demonstrates a viable approach toward solving the problem but contains some major pitfalls, errors/flaws or limitations.
- **D) Unsatisfactory (=25%)**  
Critical elements of the solution are missing or significantly flawed.  
Solution does not demonstrate sufficient understanding of the problem and/or any reasonable directions to solve the problem.
- **F) Not attempted (=0%)**  
No solution provided.

The points on a given homework question will be equal to the percentage assigned (given by the letter grades shown above) multiplied by the maximum number of possible points worth for that question. For example, if a question is worth 6 points and the answer is awarded a *B* grade, then that implies 4.5 points out of 6.